CREATE CI PIPELINE THAT WILL DEPLOY AND CONFIGURE THESE VMs USING TERRAFORM AND ANSIBLE

→ ACCORDING FOLLOWING REQUIREMENTS

• Deploy 2 virtual machines using terraform.

first vm on Amazon linux, hostname: c8.local

second vm on ubuntu 21.04, hostname: u21.local

- 2. As a result of terraform execution, dynamically create inventory for ansible c8.local should be in the frontend group u21.local should be in the backend group
 - 3. Create ansible playbook for c8.local and u21.local

for linux OS playbook should apply the following changes

selinux: disable

firewalld: disable

- 4.for frontend playbook group should install and configure nginx nginx configuration should do proxying from port 80 on port 19999 to the backend group
 - 5.for the backend group, the playbook must install the Netdata application from the

official repositories and run it on port 19999.

- CREATE A ANSIBLE SERVER ON AWS Launch instances
- → LOGIN TO SERVER THROUGH SSH AND INSTALL ANSIBLE THROUGH VISUAL STUDIO

PS C:\Users\syedf\Desktop> ssh -i "ohio.pem" ec2-user@ec2-18-119-103-199.us-east-2.compute.amazonaws.com

→Install Ansible and Terraform on Your EC2 Instance

sudo yum update -y
sudo amazon-linux-extras install ansible2 -y
ansible --version
sudo yum install -y yum-utils shadow-utils
sudo yum-config-manager --add-repo
https://rpm.releases.hashicorp.com/AmazonLinux/hashicorp.repo
sudo yum -y install terraform

→ CHECK FOR VERSION

[ec2-user@ip-172-31-11-120 ~]\$ terraform --version Terraform v1.10.2 on linux amd64

[ec2-user@ip-172-31-11-120 ~]\$ terraform --version Terraform v1.10.2 on linux_amd64

→ DOWNLOAD AWS CLI AND CONFIGURE IT

[ec2-user@ip-172-31-11-120 ~]\$ aws configure

AWS Access Key ID [None]: AKIAVVZPCK2JLYC6PJME

AWS Secret Access Key [None]: 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

Default region name [None]: us-east-2

Default output format [None]: json

→ Generate an SSH Key Pair

We'll generate a new SSH key pair to use for your Terraform-managed instances.

ssh-keygen -t rsa -b 2048 -f ~/.ssh/pipeline_key -N ""

→ CREATE A TERRAFORM DIRECTORY AND main.tf FILE

mkdir terraform_project cd terraform_project vi main.tf

→ ADD THE FOLLOWING TEMPLATE WHICH WILL CREATE TWO INSTANCE ONE WITH LINUX IMAGE AND ONE WITH UBUNTU IMAGE AND GIVE THE PUBLIC IP AS OUTPUT

```
provider "aws" {
region = "us-east-2"
resource "aws_key_pair" "pipeline" {
 key name = "pipeline"
public key = file("~/.ssh/pipeline key.pub")
resource "aws_instance" "c8_linux" {
 ami
            = "ami-088d38b423bff245f"
 instance_type = "t2.micro"
 key name = aws key pair.pipeline.key name
 associate_public_ip_address = true
 tags = {
 Name = "c8 linux"
 user data = <<-EOF
 #!/bin/bash
 hostnamectl set-hostname c8.local
EOF
}
resource "aws instance" "u21 ubuntu" {
        = "ami-00eb69d236edcfaf8"
 ami
instance_type = "t2.micro"
 key name = aws key pair.pipeline.key name
 associate_public_ip_address = true
 tags = {
 Name = "u21 ubuntu"
 user data = <<-EOF
 #!/bin/bash
 hostnamectl set-hostname u21.local
EOF
output "amazon linux public ip" {
value = aws instance.c8 linux.public ip
output "ubuntu_public_ip" {
 value = aws instance.u21 ubuntu.public ip
```

- → RUN THE FOLLOWING COMMANDS
- # terraform init
- # terraform plan
- # terraform apply

```
Outputs:
amazon_linux_public_ip = "3.147.127.28"
ubuntu_public_ip = "3.145.33.66"
```

→ checking on aws console if instance is ok!

ansible	i-0a28c286c0b36	⊗ Running ♥ ♥	t2.micro	⊘ 2/2 checks passec View alarms +	us-east-2a	ec2-18-119-103
c8_linux	i-0519ff264e9f9c	⊗ Running	t2.micro		us-east-2a	ec2-3-147-127-
u21_ubuntu	i-06382f5533664	⊘ Running ② ○	t2.micro		us-east-2a	ec2-3-145-33-6

- → NOW GET THE OUTPUT WHICH ARE PUBLIC IP'S OF THE CREATED SERVER AND ADD TO ANSIBLE INVENTORY
- # terraform output <copy the ip's>
- #cd/etc
- # sudo chown -R ec2-user:ec2-user ansible
- # vi /etc/ansible/hosts

[frontend]

Public ip ansible user=ec2-user ansible ssh private key file=~/.ssh/pipeline key

[backend]

Public_ip ansible_user=ubuntu ansible_ssh_private_key_file=~/.ssh/pipeline_key

cat /etc/ansible/hosts

[frontend]

3.147.127.28 ansible_user=ec2-user ansible_ssh_private_key_file=~/.ssh/pipeline_key

[backend]

3.145.33.66 ansible user=ubuntu ansible ssh private key file=~/.ssh/pipeline key

→ Ensure SELinux and Firewalld are Disabled

→ MAKE SURE CONNECTIVITY IS ESTABLISHED

```
[ec2-user@ip-172-31-11-120 ~]$ ansible all -m ping
3.145.33.66 | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python3"
    },
    "changed": false,
    "ping": "pong"
}
[WARNING]: Platform linux on host 3.147.127.28 is using the discovered Python interpreter at /usr/bin/python, but future installation of another Python interpreter could change this. See https://docs.ansible.com/ansible/2.9/reference_appendices/interpreter_discovery.html for more information.
3.147.127.28 | SUCCESS => {
        "ansible_facts": {
            "discovered_interpreter_python": "/usr/bin/python"
        },
        "changed": false,
        "ping": "pong"
}
```

vi playbook.yml

```
- name: Configure Amazon Linux and Ubuntu servers
 hosts: all
 become: yes
 tasks:
  - name: Disable SELinux on Amazon Linux
   when: ansible distribution == "Amazon"
   command: setenforce 0
   ignore_errors: yes
  - name: Disable firewalld on Amazon Linux
   when: ansible distribution == "Amazon"
   service:
    name: firewalld
    state: stopped
    enabled: no
  - name: Disable firewalld on Ubuntu
   when: ansible distribution == "Ubuntu"
   service:
    name: ufw
    state: stopped
    enabled: no
```

ansible-playbook playbook.yml –syntax-check

```
[ec2-user@ip-172-31-11-120 ~]$ ansible-playbook playbook.yml --syntax-check playbook.yml
```

→Install and Configure Nginx on the Frontend (Amazon Linux)

→ CREATE TEMPLATE THAT ENSURES NGINX PROXIES REQUESTS FROM PORT 80 TO PORT 19999 ON THE BACKEND SERVER.

```
# mkdir template
# cd template
# vi nginx.conf.j2
```

```
worker processes 1;
events {
  worker connections 1024;
http {
  upstream backend_group {
    {% for host in groups['backend'] %}
    server {{ host }}:19999;
    {% endfor %}
  }
  server {
    listen 19999;
    location / {
      proxy pass http://backend group;
      proxy set header Host $host;
      proxy set header X-Real-IP $remote addr;
      proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
      proxy set header X-Forwarded-Proto $scheme;
    }
  }
```

→ CREATE A PLAYBOOK TO INSTALL NGINX AND PROXY ITS PORT

vi nginx.yml

- name: Configure frontend servers (Amazon Linux) hosts: frontend become: yes tasks: - name: Add the official Nginx repository get url: url: http://nginx.org/packages/centos/7/x86_64/RPMS/nginx-1.22.1-1.el7.ngx.x86_64.rpm dest: /opt/nginx.rpm - name: Install Nginx from the downloaded RPM yum: name: /opt/nginx.rpm state: present - name: Configure nginx proxy template: src: templates/nginx.conf.j2 dest: /etc/nginx/nginx.conf - name: Start and enable nginx service: name: nginx state: started enabled: yes

ansible-playbook nginx.yml

→Install and Run Netdata on the Backend (Ubuntu)

vi netdata.yml

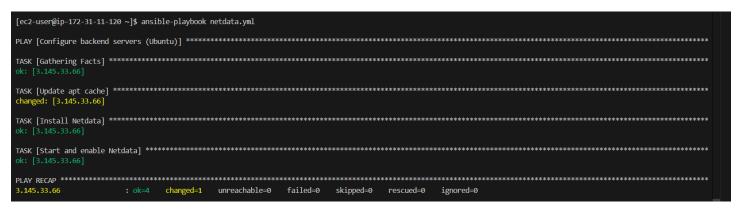
--- name: Configure backend servers (Ubuntu)
hosts: backend
become: yes

tasks:
- name: Update apt cache
apt:
 update_cache: yes

- name: Install Netdata
apt:
 name: netdata
 state: present

- name: Start and enable Netdata
service:
 name: netdata
 state: started
 enabled: yes

ansible-playbook netdata.yml



*Amazon Linux (Frontend Server): Nginx is installed, configured, and running to proxy requests from port 80 to port 19999 on the backend server.

Ubuntu (Backend Server): Netdata is installed and running on port 19999.*

< INSTANCE RESTARTED SO THE PUBLIC IP GOT CHANGED>

→ CHECKING ON BROWSER GET THE PUBLIC IP OF NGINX AND SEARCH BY PORT

19999←

Public IPv4 address

□ 3.149.237.241 | open address 🖸

< http://3.149.237.241:19999 >

